



Resolving 'embedding $\neq$ speech': the adapter trick

Standard HuBERT usage



*[X] conv extractor **BYPASSED***

CAARMA discriminator usage



Same Transformer stack, different input pipeline